

Access to Justice in the Age of AI: Evidence from U.S. Federal Courts*

Anand V. Shah

Joshua Y. Levy

Massachusetts Institute of Technology

University of Southern California

March 2026

Abstract

This paper studies how generative AI has reshaped entry into the federal civil court system. Drawing on administrative records covering more than 4.5 million non-prisoner federal civil court cases from FY2005-FY2026 and 46 million PACER docket entries matched to those cases, we document three sets of findings. First, the number of *pro se* cases—or self-represented cases without a lawyer—is increasing dramatically, rising from a long-term steady-state average of 11% to 16.8% in FY2025. This increase is concentrated in case types characterized by formulaic document production and absent from more complex, attorney-intensive categories. Second, we argue these cases are placing larger burden on federal district courts. *Pro se* cases are not terminating faster, and this combined with the increased case numbers suggests more cases for judges to process. Moreover, intra-case activity is up, with the number of motions filed within the first 90 days of a case having increased by > 40% from pre-AI means. Third, we directly validate that these increases are from AI. Using a random sample of 1,600 complaints drawn from an 8-year period (2019-2026), we find that a large and growing share of complaints are flagging positive for AI-generated text, from essentially zero in the pre-AI period to more than 18% in 2026. Together, these results suggest that the advent of large language models has increased access to justice for the common American. Whether existing institutions are equipped to handle the associated increased demand for court services remains an open question.

Keywords: access to justice, generative AI, pro se litigation, case duration, settlement, legal technology

*We'd like to thank Bao Kham Chau, Sara Jain, Jordi Weinstein, and Parker Whitfill for comments. We are also grateful to Pangram Labs for access and support. All errors are our own. JEL Codes: K40, J24, O33, L86.

It's easy with these tools to churn out hundreds of thousands of works in the time that a human author would produce maybe one or two... what we have is a room of screaming toddlers, and we can't hear the people we're trying to listen to.

— Feb. 21, 2023, Neil Clarke, editor of *Clarkesworld* on a flood of LLM-generated submissions

1. Introduction

For those in the business of science fiction, the five time Hugo Award-winning magazine *Clarkesworld* is well-known as one of the most prestigious literary venues. Yet, in February 2023, the magazine faced a crisis. Inundated with quintuple the typical submission count, the entire magazine chose to temporarily shut down, unable to scale their editorial capacity to the increased submissions attributed to unprecedentedly powerful LLMs. Three years and many models later, we find ourselves still asking how best to adapt to the decreased cost of generating adversarially passable information through our evolved filtering mechanisms.

Yet, while magazines can hire more editors, there is no easy margin along which to ‘buy’ extra judge capacity. Already case backlog is becoming a persistent feature of the federal judicial system, there is no coming influx of judges to supply additional capacity,¹ and federal courts in the United States cannot wholesale decline to hear cases (Menell and Vacca, 2021; Administrative Office of the U.S. Courts, 2024). If generative AI dramatically lowers the cost of self-represented litigation, the resulting surge in filings could overwhelm a system that depends on human judgment at every stage of adjudication. This paper examines whether that glut has arrived and, if so, how the courts are digesting it.

A reasonable way to measure the cost of entry into courts is along people choosing to file *pro se*. The right to represent oneself in legal proceedings, or *pro se* litigation, is among the oldest statutory rights in American law, older even than the Bill of Rights.²³ For most of our

¹Only an act of Congress can expand the number of Article III judges, and the number of Article III judge appointments has grown only modestly for decades.

²The right to appear *pro se* in federal civil cases was first enacted as Section 35 of the Judiciary Act of 1789, signed into law on September 24, 1789—one day before the Bill of Rights was proposed to the states.

³Will Hunting: “I am afforded the right to speak in my own defence, sir, by the Constitution of the United States! This is the same document which guarantees my liberty. And liberty, in case you’ve forgotten, is a soul’s right to breathe. And when it cannot take a long breath, laws are girded too tight. Without liberty, man is a syncope.” Judge: “Man is a what?” Will Hunting: “Ibid, Your Honour.” – *Good Will Hunting*

courts’ history, this right has been exercised at a ‘remarkably unremarkable’ steady state. Gough and Taylor Poppe (2020) examine federal civil filings from 1999 to 2018 and document a civil litigation *pro se* representation rate of approximately 11%—stable across two decades and notably unperturbed by the Great Recession, the Affordable Care Act litigation wave, or the surge in immigration cases under the first Trump administration.

This stability seems to reflect a structural barrier: for most people, self-representation is prohibitively hard. Filing a federal civil complaint requires identifying the correct jurisdictional basis, pleading sufficient facts to survive a motion to dismiss, and navigating procedural requirements that vary by court and case type.⁴

The widespread, public diffusion of powerful LLMs changes that calculus. Without a law degree and at *de minimis* cost, any person with an internet connection can not only obtain interactive, case-specific legal guidance—drafting complaints, identifying statutes, navigating procedure—but also generate passable legal documents, particularly so after the release of GPT-4 in March 2023.⁵ Whether these technological advances have been sufficient to materially alter filing behavior is an important empirical question for the future of our courts. We study it using administrative records covering roughly 4.6 million non-prisoner federal civil cases filed from fiscal years 2005-2026. We draw on two complementary datasets: filing-based counts from the Administrative Office of the U.S. Courts and the Federal Judicial Center Integrated Database (IDB), which tracks filings, dispositions, and durations at the case level. We supplement the IDB with 46 million PACER docket entries that enable us look inside cases to measure motion activity.

We document three sets of findings.

First, the national non-prisoner *pro se* filing share rose sharply from its approximately 11% historical steady state to 16.8% in fiscal year 2025, a gain that has no precedent in 25 years of

⁴This is despite the “liberal construction” standard established in *Haines v. Kerner* (1972), in which federal courts try to read *pro se* filings charitably, often overlooking technical defects.

⁵At the time of its release, the two most widely cited legal-reasoning evaluations for large language models were the Uniform Bar Examination—the standardized test used to admit new lawyers in 46 U.S. jurisdictions—and LegalBench, a 162-task academic benchmark spanning issue spotting, rule application, legal interpretation, and rhetorical analysis. Katz et al. (2024) report that GPT-4 scored approximately 297 on a full simulated UBE, above the passing threshold in every jurisdiction and at roughly the 75th percentile of February 2022 test-takers. Guha et al. (2023) find that GPT-4 was the best-performing commercial model on LegalBench at release, with category averages of 83% on issue spotting, 90% on rule conclusion, 75% on interpretation, and 79% on rhetorical understanding.

administrative records. Moreover, that rise is not uniform across case types. It is concentrated in categories where the dominant task is producing a well-structured factual narrative—civil rights complaints, consumer credit disputes, foreclosure proceedings—and is essentially absent in categories that require sustained expertise, such as patent infringement and securities fraud.

Second, cases are not resolving any faster and analysis of case disposition suggests little change in the quality of cases brought by *pro se* litigants. However, more is happening within each case. Since the advent of capable LLMs, courts have experienced a 209% increase in the number of motions being filed in *pro se* cases in the first 90 days. A motion is a claim on the judge’s time: it requests some judge intervention or expertise along the presented facts of the case. These results bundled suggest increased court burden in the post-AI period.

Third, we directly test that AI use is increasing in the post-period by applying an AI-content detector to a random sample of 1,600 federal civil complaints drawn from CourtListener’s RECAP archive. The AI-content detector records only one false positive (out of 800 queries) in the pre-period. We find that the share of *pro se* complaints containing AI-generated text has been rising monotonically from 1.0% in 2023 to 18.0% in early 2026.

Finally, an important caveat is this paper is necessarily descriptive, given every internet user was given access to each LLM model at once, and that LLM use has diffused quite widely (see Figure 5). Moreover, differences in penetration are hard to justify without fine-grained usage data, given LLMs have seen the fastest adoption curve in the history of technology.⁶ The only counterfactual available is the pre-AI time series in each unit, and so the identification argument throughout this paper is at best temporal stability: the twenty years of *pro se* filing behavior that precede GPT-4 look the same as the ten years that precede it, or the five (absent COVID-19), so the post-2022 break is unlikely to be a continuation of a secular trend. The visual tests in Figures 1, 4, and 5 are meant to make this stability easy to see and to let the reader judge for themselves how unusual the post-LLM break is relative to earlier variation. We do not claim to identify a causal effect of GPT-4 on *pro se* filing, only that the observed time series is difficult to rationalize without generative AI playing a role, and that the cross-sectional pattern matches the simplest prediction such a role would generate.

The paper proceeds as follows. Section 2 reviews the related literature and describes the data.

⁶ff

Section 3 documents the headline *pro se* results, characterizing who is filing and in what case types. Section 4 examines the outcomes of these filings, including durations, dispositions, and intra-case docket activity. Section 5 reports the results of an AI-content detector applied directly to the complaint documents. Section 7 concludes.

2. Background, Data, and Methods

This paper contributes to three literatures. First, a large body of scholarship documents the barriers that the cost of legal representation creates for low-income individuals (Rhode, 2004; Sandefur, 2019; Hadfield, 2016; Legal Services Corporation, 2022), and experimental evidence shows that legal assistance materially improves outcomes for those who receive it (Greiner et al., 2013). We provide what is, to our knowledge, the first large-scale quantitative documentation that a general-purpose AI technology has altered the rate at which individuals choose to self-represent in federal court.

Second, randomized and quasi-experimental studies document substantial productivity gains from large language models in writing (Noy and Zhang, 2023), customer service (Brynjolfsson et al., 2023), consulting (Dell’Acqua et al., 2023), and job search (Wiles et al., 2025), with gains often concentrated among lower-ability workers. Our setting differs in that the beneficiaries are not necessarily professional workers but individuals accessing a formal institution.

Third, models of litigation selection (Priest and Klein, 1984; Shavell, 1982; Waldfogel, 1995) predict that improvements in the information environment increase settlement rates, a prediction that Siegelman and Waldfogel (1999) confirm empirically. The settlement and adjudication patterns we document are plausibly consistent with this mechanism.

2.1. The Federal Civil Litigation Process

Federal civil cases are governed by the Federal Rules of Civil Procedure, which impose uniform pleading, discovery, and motion practice requirements across all 94 U.S. district courts (our data universe). A case begins with the filing of a complaint, which must identify the court’s basis for subject-matter jurisdiction, the legal theory supporting the claim, and the factual basis for relief. Cases are classified at filing by Nature of Suit (NOS) code, a four-digit taxonomy that distinguishes among approximately 80 substantive legal categories ranging

from civil rights and consumer credit to patent infringement and securities fraud.

The *pro se* litigant occupies a formally protected but practically disadvantaged position in this system. Courts are required to construe *pro se* filings liberally and many districts maintain self-help centers offering procedural guidance. Nevertheless, *pro se* litigants consistently achieve worse case outcomes than represented parties, even after conditioning on case type and court (Greiner et al., 2013; Steinberg, 2023). The gap reflects not only deficiencies in legal knowledge (as *pro se* cases are on average lower quality) but also disparities in the capacity to navigate procedural requirements, respond to motions, and engage in discovery—a result that we document in Section 4. [avs: refer to figure when we have it]

2.2. Data Sources

2.2.1. FJC Integrated Database (Duration and Disposition Analysis)

For the analysis of trends in case filing, case duration, and case outcomes, we use the Federal Judicial Center Integrated Database (IDB), which records case-level information on all civil cases terminated in all U.S. district courts (Federal Judicial Center, 2025). Our data spans data files from FY2005-2025, which we combine into a single deduplicated dataset indexed by unique case identifiers.

The resulting sample, after excluding prisoner NOS codes (510, 530, 535, 540, 550, 555, 560) and MDL direct-filed cases (ORIGIN code 13), contains approximately 3.7 million cases with valid filing dates, termination dates, and positive recorded durations. For each case we observe: the NOS code, the district, the filing date, the termination date, the representation status (PROSE field, coded 0 for represented, 1 for plaintiff *pro se*, 2 for defendant *pro se*, 3 for both parties *pro se*), the disposition code (DISP), and the judgment outcome (JUDGMENT).

2.2.2. AOUSC C-13 Table

For filings data, we augment the IDB data set with the C-13 table published by the Administrative Office of U.S. Courts (AOUSC), which records civil case filings by district, year, and representation status. This dataset includes only district filing data aggregated by year.

In both the IDB and C-13 datasets, we exclude two populations of case filers. First, we exclude cases from prisoner filings (cases arising under 28 U.S.C. §§ 2241, 2254, 2255, and

related statutes), which constitute a qualitatively distinct population responding to different institutional incentives and under different *pro se* requirements. We also exclude organized mass-litigation events (in C-13, this is; in IDB, this is...). We do so as such organized mass-filing events (often related to product recalls) depart from the qualitative features that characterize typical civil *pro se* litigation, thereby obfuscating from the interesting economic and social phenomena at play.

2.2.3. PACER DATA from MC!

We also access metadata for 46 million PACER docket entries and match them to IDB cases to characterize motion activity within each case. Federal district courts differ in how they publish their docket data. For the 54 district courts (of 94 total) that provide full PACER feed coverage, we achieve a 99.8% match rate to the underlying case data.

2.3. Methods

2.3.1. Correcting for right-censoring and termination bias

Our analysis concerns patterns associated with both case filings and case outcomes. However, almost tautologically, we can only observe outcomes for cases that have terminated. So, any analysis of outcomes must account for selection from right-censoring. To illustrate, imagine there are only two equally likely cases, cases that resolve in 6 months c_A and cases that resolve in 18 months c_B . Suppose further that we can observe data only through December 31, 2025. For cases filed sufficiently far in the past, we would observe that cases have an average duration of ≈ 1 year (including both c_A and c_B terminations), whereas as we approach the right-edge of the time series we would observe an average duration of ≈ 6 months (including only c_A terminations), and thereby incorrectly assume that recent cases resolve doubly fast! Indeed, if we compute any outcome measure on the right-edge of the sample and compare it to the left-edge to view the difference in that measure over time, we'll be polluting that comparison by also implicitly measuring the difference between $\{c_A\}$ and $\{c_A, c_B\}$ cases. This is the termination-bias problem.

We implement a horizon-specific, nonparametric estimator to account for right-censoring of outcomes by restricting the risk set at each duration d to cases whose observation window (the number of days between filing date and the database edge date of December 31, 2025) is

at least d days. At each point d , the resolution rate $\hat{F}(d)$ is then:

$$\hat{F}(d) = \frac{\#\{i : \text{obs_window}_i \geq d, \text{duration}_i \leq d\}}{\#\{i : \text{obs_window}_i \geq d\}}$$

Returning to the original example concerning case termination rates, suppose we observe N cases on December 31, 2025, only a subset $n < N$ of which have terminated. For a duration d , discard cases that have not been observed for at least d days. For instance, if we choose $d = 180$, we discard all cases (terminated or otherwise) that were filed after July 4, 2025. By decreasing the duration parameter, our estimator trades off allowing more recent cases to contribute to our estimate of the termination rate $\hat{F}(d)$ at the cost of restricting the temporal horizon over which we can make a statistically precise claim (the true case termination rate). By contrast, larger values of d permit longer-run case resolution, but require dropping more recent filings, thereby reducing sample size and making estimates less precise. Put differently, choosing d involves a tradeoff between precision and the temporal horizon of the outcome being measured.

For the purposes of our analysis, we face an additional tradeoff. In order to comment on the nature of all cases (as opposed to just those that resolve quickly) we would like to pick a large duration parameter, d . Picking large values for d mechanically truncates more data from the end of our time-series, however, the sample which would capture the effects of AI diffusion on outcomes. That is, large observations windows come at the cost of losing post-AI period data, by construction. In light of this tradeoff, we present analysis for four different observation windows: $d \in \{90, 180, 270, 365\}$. We apply such observation-window corrections to all of our outcome measures.

[avs: we will nail this tomorrow] [jyl: I think I have cogent/readable explanation now]

3. Who is filing, and in what cases?

We document a steep rise in the number and share of *pro se* filings in federal civil court since FY2022, driven primarily by a growing number of self-represented plaintiffs. We also find that this rise concentrates in types of cases where document production constitutes the primary cost of filing, and that this rise has taken place almost uniformly across states.

3.1. The national *pro se* share and count

Figure 1 illustrates this rise using both of our primary datasets. The *pro se* share of non-prisoner filings hovered around 11% from FY2005 through FY2022, a level Gough and Taylor Poppe (2020) described as “remarkably unremarkable.” In FY2025 it reaches 16.8%. The change is continuous and steep rather than a one-year anomaly: shares rise from 11.6% in FY2022 to 11.4% in FY2023 (flat), 14.5% in FY2024, and 16.8% in FY2025. Represented filings do not fall. [jyl: These contemporaneous trends can be interpreted in two ways: First, it is possible that the additional *pro se* cases are new entrants to the federal civil courts system who would not have otherwise filed a case represented by professional counsel and the diffusion of popular LLMs in late 2022 and early 2023 did not increase the number of cases being brought by represented parties. Alternatively, it is possible that the advent of these generative AI tools induced some parties to substitute away from professional legal representation and to file their lawsuits *pro se*. However, this substitution effect must have, on net, been dominated by the entry of new represented parties because we do not observe a net decrease in the number of represented cases. [EXTRA SPECULATIVE PART: Supposing that the the advances of LLMs provide a disproportionate benefit to *pro se* litigants who face relatively large cognitive and procedural fixed costs to filing a suit relative to trained lawyers, we prefer the first interpretation.]] The additional *pro se* cases are new entrants into the federal civil court system, not substitutes for cases that would otherwise have been brought with counsel.

Figure 2 plots non-prisoner civil filings in federal district court by fiscal year, separated into represented and *pro se* cases. Two facts stand out. First, represented filings are essentially constant over this period: they averaged roughly 193,000 per year between FY2005 and FY2022 and stood at 205,890 in FY2025, squarely inside the range of outcomes of the pre-AI period. Second, *pro se* filings have broken out of their long-run range. They averaged 23,210 cases per year from FY2005 to FY2022, never straying far from that level through the Great Recession, the Affordable Care Act litigation wave, or the 2017–2018 episode of immigration-related filings during the first Trump administration. Beginning in FY2023, *pro se* filings jump: 27,370 in FY2023, 31,478 in FY2024, and 41,490 in FY2025. The FY2025 count is almost double the pre-AI mean.

Using the AOUSC filing data, by FY2025, total civil litigation had increased by 18,280 cases

The national *pro se* share of civil litigation is rising rapidly




Figure 1. National non-prisoner *pro se* share, FY2005–FY2025. Twenty years of stability followed by a sharp post-2022 rise. Source: AOUSC C-13.

relative to the pre-AI mean (an 8.5% increase in counts). However, despite representing a share of all filings an order of magnitude smaller than represented cases, new *pro se* filings account for 59.3% of the growth in civil filings. [jyl: We interpret this as being consistent with the conjecture that the widespread of adoption of generative AI tools is driving increased filings by those who gain disproportionately from the new technology at the filing stage: non-lawyers.]

3.2. Plaintiffs, not defendants

A *pro se* case can arise for two very different reasons. A plaintiff may choose to file without counsel because they have a claim they want to pursue and cannot or will not incur the costs of hiring a lawyer to pursue it. A defendant may appear in court without counsel because they were sued and cannot afford a professional defense. [jyl: Because plaintiffs initiate civil legal proceedings, their choice to forgo professional representation is materially different than defendants’ choice to do the samme.]The first is entry; the second is constraint. Figure 3 decomposes the annual *pro se* case count into these roles using the IDB PROSE field, which distinguishes plaintiff-side, defendant-side, and both-sides self-representation.

The post-AI rise in self-representation is a plaintiff story. Plaintiff-side *pro se* case counts averaged 20,349 per year from FY2015 to FY2022 and reach 39,167 in 2025, nearly doubling.

Federal civil litigation has reached an all-time peak, in part driven by an increase in *pro se* filings




Figure 2. Non-prisoner federal civil filings, FY2005–FY2025. Represented filings are essentially flat; *pro se* filings step up sharply starting in FY2023. Source: AOUSC C-13.

Pro se role composition: plaintiff entry drives the post-AI rise




Figure 3. Pro se case counts by role, FY2015–FY2025. The post-AI rise loads entirely on plaintiff-side *pro se* filings. Source: FJC IDB (cv26), non-prisoner non-MDL.

Defendant-side *pro se* counts fall slightly over the same window, from 4,650 to 3,896. Both-sides *pro se* cases grow modestly, from 775 to 1,081. The increase in the total *pro se* count since FY2022 is therefore accounted for by plaintiff-side entry; defendant-side self-representation, if anything, has receded. The fact the rest of the paper tries to explain is not that people are being denied counsel. [jyl: We conclude that people are choosing to enter the federal judicial system and file complaints of their own accord at a rate without modern historical precedent.] It is that people are walking into federal court on their own accord and filing complaints, at a rate with no precedent in the past two decades.

3.3. Simple cases, not complex ones

If generative AI has lowered the cost of producing filing documents, the rise in *pro se* entry should concentrate in case types where filing-stage document production represents the dominant cost in initiating and pursuing a lawsuit. [jyl: For civil rights or consumer credit cases, which often resolve at or shortly after the pleading stage, AI can complete at *de minimis* cost many of the tasks that a lawyer would complete for their client, but are extremely demanding for even educated laymen] For civil rights or consumer credit cases, which often resolve at or shortly after the pleading stage, AI can substitute for much of the work a lawyer would otherwise do. For patent infringement or securities fraud, which require

sustained discovery, adversarial motion practice, and expert testimony, AI can substitute for comparatively little. So, does the post-AI rise in *pro se* litigation actually load on the case types where filing-stage help matters most?

We use the pre-AI *pro se* share of each Nature of Suit (NOS) category as a revealed-preference measure of how amenable the case type is to self-representation. NOS categories whose pre-2022 *pro se* shares were high are categories where many litigants, even before GPT-4, had concluded that self-representation was feasible.⁷ NOS categories whose pre-2022 *pro se* shares were low are categories where self-representation was rare, presumably because the case type required more than filing-stage work. We split the 86 non-prisoner NOS codes at the case-weighted median of pre-AI *pro se* share, yielding 48 “simple” and 38 “complex” categories. The category contains, among others, civil rights (NOS 440, 42% pre-AI *pro se*), foreclosure (220, 31%), and consumer credit (480, 9%, the marginal code). The complex bucket contains patent (830), securities (850), ERISA (791), and medical malpractice (835), which all exhibited a pre-AI *pro se* share below 5%.

Figure 4 plots the *pro se* share for each category by fiscal year. The two series track each other from FY2015 through FY2022 and then diverge. Simple-category lawsuits’ *pro se* share rises from roughly 22% (pre-AI mean) to 30.5% in FY2025. The analog for complex cases rises from 4.9% to 7.1%. In raw terms, the entire post-2022 *pro se* rise at the national level loads into the 48 NOS categories that were already above-median amenable to *pro se* entry before the AI era. It is almost entirely absent from the 38 categories that were not. [jyl: Our finding strongly coheres to the simplest version of the AI-as-filing-assistant hypothesis.] This is what the simplest version of the AI-as-filing-assistant hypothesis predicts, and it is what we find.

Our above/below-median partition is constructed using exclusively pre-AI data: no post-FY2022 information enters the classification procedure. That a pre-AI partition successfully separates the post-AI *pro se* surge is not a property that would hold if the rise were driven by a secular trend in the population-level preference for self-representation, demographic shifts, or a change in case composition driven by policy or extraordinary events. Rather, it is a property that can be interpreted as consistent with a cost shock that acts selectively through

⁷This may have been due to *ex ante* variability in the templatability of certain types of suits, the availability of specialized legal-aid or self-help clinics, and other informational resources.

the filing-stage component of legal work.




Figure 4. *Pro se* share by fiscal year, split at the pre-AI case-weighted median of *pro se* share. The simple bucket rises sharply post-2022; the complex bucket barely moves. Source: FJC IDB (cv26), non-prisoner non-MDL.

3.4. Geographic spread

The third question arises about where these cases are being filed. If the *pro se* rise reflects something local—a single district’s policy change, a courthouse that became especially welcoming to self-represented filers, or a state bar association or an area law school’s initiative—we would expect the increase to be concentrated in a few places. If instead it reflects a technology shock available everywhere at the same time, we would expect the increase to show up almost everywhere at once. [avs: hex plot is missing md and idaho – look into it?]

Figure 5 aggregates district-level *pro se* filing counts to the state level and plots FY2025 relative to each state’s peak pre-AI year. For a state to appear “teal” in this figure, FY2025 *pro se* filings must exceed the single highest pre-AI year in that state’s own history. All but six states clear this bar.⁸ The rise is not driven by a handful of outlier districts. It is broad, immediate, and geographically pervasive in a way that is difficult to reconcile with any explanation local to a specific jurisdiction.

⁸A notable outlier is Vermont, which reflects strategic “forum shopping” by a specific group of *pro se* immigration litigants rather than a broad trends; excluding Vermont leaves every result in the paper unchanged and we provide a more detailed case study of the District of Vermont in Appendix A.

Relative to peak pre-AI *pro se* activity, FY2025 *pro se* filings were higher almost uniformly across the country




Figure 5. FY2025 non-prisoner *pro se* filing count relative to each state’s peak pre-AI year (highest annual count across FY2005–FY2022). Darker cells indicate states where FY2025 exceeds the pre-AI peak. Source: FJC IDB (cv26).

4. What are courts doing with these cases?

The previous section documented a large, rapid, and geographically pervasive increase in federal *pro se* civil filings concentrated in simple case types. What happens to those cases once they enter the system? If AI is helping self-represented plaintiffs draft complaints but nothing else, the additional cases should resolve at around the same pace as pre-AI cases and leave little trace on the judicial system. If AI is helping *pro se* parties sustain litigation further into the process, durations should shift, dispositions should change, and motion activity should rise.

Three caveats apply throughout this section. First, all of our outcome measures (within-case activity, case durations, dispositions, and judgments) are subject to right-censoring: we observe outcomes only for cases that have terminated as of December 31, 2025, so cases that would take a long time to resolve are under-represented in the most recent cohorts. To correct for this, we apply observation-window corrections throughout, following the estimator in Section 2.2. Second, the COVID-19 pandemic is a large confounder for any outcome measured on filing cohorts between March 2020 and late 2021; we read the post-AI series against the pre-COVID baseline (FY2015–FY2019) and flag COVID artifacts we assess them to have

material implications for our analysis. Third, the PACER docket data used for intra-case activity measures covers the 54 federal districts with full-feed PACER access, a subset that accounts for roughly 80% of non-prisoner federal civil cases.

4.1. Case durations are unchanged

The simplest outcome to measure is how long cases take to resolve. If the new *pro se* filings are meritless and being dismissed on the papers, average case duration should fall. If they are sustaining litigation further into the process, it should rise. Figure 6 plots the share of cases resolved by 90, 180, and 365 days from filing, computed at the monthly filing cohort level with observation-window correction.




Figure 6. Monthly cohort resolution rate at 90, 180, and 365 days from filing, FY2015–FY2025. Each dot is one filing-month cohort. All three series return to the pre-COVID band after 2022 and remain there through the end of the sample. Dashed lines mark ChatGPT (Nov 2022) and GPT-4 (Mar 2023). Source: FJC IDB (cv26), non-prisoner non-MDL, obs-window filtered.

At all three horizons, the post-AI resolution rate sits inside the pre-2020 noise band. Cases filed in FY2023–FY2025 resolve no faster or slower than cases filed in FY2015–FY2019. The COVID period produces a transient dip in resolution rates that persists through roughly mid-2021 and recovers afterward; the post-AI period is a return to the long-run baseline, not

a departure from it. By the simplest measure of court throughput—the time from filing to termination—AI has left no visible mark.

4.2. But docket activity inside cases has risen sharply

Case duration statistics may generally summarize the arc of cases but say nothing about what happens along the way. A second class of measures looks *inside* cases: motion practice, hearings, and other docket activity. These are the measures that capture the burden a case places on the court while it runs its course, and they are where AI has left its clearest mark.

Figure 7 plots the number of motions filed in the first 90 days after a case is filed, quarterly, separately for *pro se* and represented cases. The left panel plots the the total number of motions filed in each court; the right panel presents the number of motions per case.

Per-case intensity is up for both groups, and up more for *pro se* cases. *Pro se* motions per case rose from 1.34 in the pre-AI period (2015–Q3 2022) to 1.53 post-AI, a 14% increase. Represented motions per case rose from 1.16 to 1.26, an 8% increase. Because the *pro se* case count is also rising, the total volume of *pro se* motion activity per district has approximately doubled. A typical full-feed district saw 82 *pro se* motions per quarter in the first 90 days of newly-filed cases before the advent of widely available AI, 145 after, and 252 in 2025Q3, the most recent fully observable quarter. At end-line, this is number is more than triple the pre-AI baseline.

4.2.1. Symmetric, not one-sided


[jyl: Can redo motion burden figure to add the represented motions (dashed line); could also do a demeaned thing] Notice that the increase is not one-sided. Represented plaintiffs also file more motions per case than they did before: 8% more in the first 90 days. Attorneys are using LLMs too, and the same cost shock that made it cheaper for a *pro se* plaintiff to draft a motion made it cheaper for a represented plaintiff’s attorney to draft one. The notable pattern is not that one side is reacting to the other. It is that both sides are filing more motions at once, so the total burden a case places on the court during its first 90 days is rising symmetrically on both sides of the bar, not just on the *pro se* side of it.

[avs: pro se vs represented, same as old figure but just both overlaid, but in shares]

Early motion burden has risen in both pro se and represented federal civil cases


Pro se 90-day motions per court-quarter

Motions per court




Pro se 90-day motions per case

Motions per case



Represented 90-day motions per court-quarter

Motions per court



Represented 90-day motions per case

Motions per case



Note: All available districts, filings from 2015-01-01, full-feed courts only; docket motions within 90 days of filing (160,658 pro se cases; 967,237 represented cases).

Source: FJC IDB × CourtListener RECAP motion feed (canonical compare series).

Figure 7. Motions filed within 90 days of complaint, quarterly, on the 54 full-feed districts. Left: motions per case, by representation. Right: total *pro se* motions per district per quarter. Source: Matt Clark PACER docket entries + FJC IDB (cv26).

4.2.2. Concentrated in simple case types

If AI is the driver, the motion-burden rise should load on the same case types where the *pro se* filing rise loaded—the simple NOS categories. Figure 8 repeats the motion-per-case and motion-per-district analysis separately for the simple and complex buckets, using the same pre-AI PS-share partition as Section 3.3.




Figure 8. 90-day motion burden on the 54 full-feed districts, by simple vs. complex NOS partition. Left: *pro se* motions per case. Right: *pro se* motions per district per quarter. Source: Matt Clark PACER docket entries + FJC IDB (cv26). Partition defined in Section 3.3.

The concentration is sharp. In simple NOS cases, *pro se* motions per case rose from 1.34

pre-AI to 1.64 post-AI (a 22% increase), and total motions per district rose from 67 to 125 (85%). In complex NOS cases, motions per case are actually lower post-AI than pre-AI (1.34 to 1.11), and motions per district rose from 16 to 22 (36%)—a smaller rise driven mostly by the modest complex-bucket filing increase rather than by any change in intra-case intensity. The additional docket activity documented in the previous subsections is almost entirely a simple-NOS phenomenon. It tracks the same partition as the *pro se* filing rise, which is consistent with a common underlying cause.

By the outcome measures a court administrator would check first—durations, dispositions—AI has left no visible trace on the federal civil courts. By the measures of what happens inside cases, it has. The pattern is consistent with a technology that lowered the cost of producing legal documents: more documents per case, more cases per court, and no change in how any given document is adjudicated once filed. To our knowledge, this is the first large-scale documentation of rising federal court docket activity attributable to generative AI. The downstream consequences of that rising burden—backlog, delay, congestion—are only beginning to be felt, and we return to them in Section 7.

4.3. Case outcomes are unchanged

[avs: Joshua disposition heterogeneity analysis: pro se win/loss/settled/judge-dismissed buckets over time. Expected finding: no large shift in the shares, i.e., new pro se cases are not disproportionately winning or being dismissed relative to the historical mix. Include plot + two to three sentences quantifying the pre/post shares.]

5. Is AI-generated text showing up in the documents themselves?

The results so far are consistent with AI being used by self-represented litigants, but they fall short of being dispositive. The filing increase, the concentration in simple case types, the geographic spread, and the rise in motion activity are patterns we would expect if AI had lowered the cost of legal document production. They could also, in principle, reflect some other shock: a demographic shift, a legal-aid funding change, or a cultural moment we have not considered. To distinguish these possibilities, we pull a random sample of complaint documents from federal courts and test whether they are generated by LLM tools.

5.1. Text detector

We use Pangram Labs, a commercial AI-text detection service, for all AI-text detection in this paper. Imas and Jabarian (2025) evaluate four leading AI-text detectors—Pangram, OriginalityAI, GPTZero, and RoBERTa—and find that Pangram is the only detector to satisfy a strict false-positive cap ($\text{FPR} \leq 0.005$) without sacrificing accuracy, achieving a false-positive rate of essentially zero and a false-negative rate between 0.5% and 3.8% depending on the threshold and model. We flag a document as a “hit” when the Pangram API determines that its contents has a strictly positive probability of having been generated by an AI tool.

5.2. Sample and method

We draw a stratified random sample of 200 federal civil complaints from each calendar year between 2019 and 2026, for a total of 1,600 documents. The sample comes from the CourtListener/RECAP archive, which mirrors PACER docket filings that contributors have requested. Our procedure proceeds as follows:

Randomization. Within each year, we randomly select a day and then a random civil case from that day (`docketNumber:cv`).

Within case. Within each case, there are many documents—these documents can be complaints, motions, or judge memos, but they can also be empty, boilerplate logistical orders, or broken text. To standardize document selection, for each case docket, we choose the *first* document on the docket `entry_number:1`. In federal civil practice, the first docket entry is almost always the initiating complaint. While sampling, for all the days where the docket is found to have a valid RECAP document, we also apply three deterministic filters.

1. First, we require a positive description match to “complaint” or “petition” (standard language on complaint documents) to exclude misfiled or mislabeled entries.
2. Second, we filter to non-prisoner NOS codes (for consistency with our sample). This constraint did not bind for any of the documents we sampled.
3. Third, we require that PDFs have at least 500 characters of text (Courtlistener applies no quality filtering of their own, and occasionally PDFs are empty beyond a boilerplate PACER header).

If the case has an available RECAP document, and if the document passes these deterministic filters, then we call the complaint ‘valid’ and retain the PDF.

Pangram. Before passing the PDF to Pangram, we strip the PDF of its automatically generated boilerplate PACER header (court name, filing date, case caption)—this boilerplate is not in the complaint as courts read them, and interestingly also generates some false positives (especially in more recent years). After stripping boilerplate, we pass the complete complaint to the Pangram API for AI-text detection.

Repeat. For each year, across 2019–2026 (four pre-period years, four post-period years), we repeat this process 200 times.

Below, we present the results of this procedure.

5.3. Results

The pre-AI period (2019–2022) yields 1 detection across 800 documents, a rate of 0.1%. The post-AI period (2023–2026) yields 66 detections across 800 documents, a rate of 8.2%. The rate rises monotonically from year to year: 1.0% in 2023, 3.5% in 2024, 10.5% in 2025, and 18.0% in early 2026. Figure 9 plots the share of flagged complaints by year.




Figure 9. Share of federal civil complaints classified as containing AI-generated text, by filing year. Sample: 200 random complaints per year, 2019–2026, drawn from CourtListener/RECAP. Error bars are Wilson 95% confidence intervals. Source: Pangram AI-text detector applied to complaint bodies.

The pre-AI baseline is essentially zero, and is conducted to calibrate Pangram’s false positive rate on this specific kind of text: legal complaints. A single detection out of 800 pre-period complaints means that the detector is not producing false positives at a rate that would confound the post-period signal. In the post-period, we find that the AI detection rate rises continuously, from near zero at the end of 2022 to 18% by early 2026, tracking the trajectory of LLM capability gains and diffusion of LLM adoption rather than any single event. By early 2026, roughly one in five complaint filings contains text that the detector classifies as AI-generated. This is likely an underestimate, because detectors tend to flag the most heavily AI-written documents and so may miss those where AI was used lightly.[\[avs: have to check the imas/jabarian paper against this last claim\]](#)

That AI use is this visible at the document level, even in a sample that includes attorney-drafted pleadings, is itself notable.⁹ Combined with the patterns in Sections 3 and 4, the evidence supports a simple reading: AI is being used to generate federal civil complaints, the use is growing quickly, and it is showing up in filing behavior and in the text of the filings themselves.

An important point of note is that, though we try to maintain a high-quality sample and randomize as completely as possible, there is one important source of bias in the sample. Namely, the universe of RECAP documents available to us is itself not the complete universe of case documents (though it is the most complete sample publicly available). This is because documents are automatically uploaded by users of the RECAP browser extension—disproportionately attorneys, journalists, and legal researchers—so that *pro se* documents specifically are likely under-represented. If *pro se* litigants are using AI at a higher rate than attorneys, that suggests that our results are a lower bound for the true prevalence of AI-text in legal cases.¹⁰

⁹Federal judges are prohibited from using AI for drafting opinions under the Judicial Conference’s 2024 interim guidance, and most district courts now require attorneys to certify whether AI was used in any filing. The Pangram detection rate is consistent with these rules being widely ignored or widely unknown. [\[jyl: Am I reading this correctly? This feels like a huge jump that we can’t claim.\]](#)

¹⁰We could only conceive of one way in which *pro se* documents might be over-represented in the post-AI period: if commercially available LLMs being used by *pro se* litigants *explicitly* encourage their users to upload AI-drafted complaints to CourtListener or to install the RECAP browser extension. In anecdotal testing, we were unable to elicit this kind of behavior in a number of widely used AI tools.

6. Avenues for Future Work

- Other legal settings – state and municipal cases where stakes are lower and so people might be more willing to delegate to the cheap option (parking tickets etc.)
- Case quality — looking inside the docket
 - Are these post-AI complaints better written?
 - Are we seeing stable durations despite more motions because people are able to file motions and counter motions faster (lets look at time between filings in PACER)
 -
- Where are we getting entry? Are people substituting away from lawyers on their second+ cases? Are all of the new plaintiff cases repeat filers?
- What is happening to the labor market for lawyers?
 - Is this causing polarization in firms? Who’s going to benefit here, big firms or small firms? – Maybe the beneficiaries are those that do the most complex litigation?
 - What are law students doing? are they pricing in these effects?
- Thinking about the docket pipeline from a task-based approach? O-ring automation?

7. Conclusion

The federal civil courts are absorbing a large, rapid, technology-driven increase in *pro se* filings. That the absorption is happening in federal court, rather than in the more accessible and lower-stakes state or municipal court systems, is itself striking. Federal civil cases require satisfying federal jurisdictional rules, surviving federal pleading standards, and navigating a procedural apparatus that is widely acknowledged to be among the most demanding in the American legal system. If generative AI is enabling self-representation in this demanding context, the effect in state and local courts—which have lower filing fees, simpler procedures, and higher case volumes—is likely larger still.

Within the federal system, the supply side of the adjustment margin is tightly constrained. Article III judges require years of training and Senate confirmation, and the authorized

number of district judgeships has grown only modestly since the last comprehensive expansion in 1990 (Menell and Vacca, 2021). Current Judicial Conference guidance prohibits federal judges from using AI to draft opinions, and most district courts require attorneys to certify whether AI was used in any filing. The supply of adjudicative capacity is therefore fixed on nearly every margin: the number of judges is fixed by Congress, the training time of new judges is fixed by the profession, and the ability of existing judges to use AI to process more cases is prohibited by policy. This paper has documented a rise in demand for adjudicative services. It is hard to see how the mismatch resolves without some form of adjustment.

Two trajectories worry us in particular. The first is an arms race between private litigants. If represented plaintiffs are already responding to rising *pro se* motion volume by filing more motions of their own, the post-AI equilibrium may feature cases where motion-per-case counts drift upward on both sides, absorbing an increasing share of judicial attention per case without resolving a correspondingly greater share of disputes. The 14% and 8% per-case increases we document are small relative to what this process could produce if it continues for years to come.

The second is asymmetric litigation against institutional defendants, particularly the federal government and large administrative agencies that cannot easily scale their litigation capacity in response. A federal agency sued by a plaintiff with AI-drafted filings faces an adversary whose marginal filing cost has fallen, while the agency’s own response cost has not. The civil rights, immigration, and benefits-appeal categories in which we find the largest *pro se* increases are disproportionately cases against federal defendants. If AI scales the rate at which such cases are filed faster than it scales the rate at which agencies can respond, the government accumulates a queue of cases that it cannot process, with consequences for administration as well as adjudication.

The structural question is whether the supply of adjudicative services can grow at all. The logic is simple: if demand for court services has risen and the policy goal is to preserve access, supply has to rise too. The margins on which it could rise are limited. One option is to relax the rules that prohibit federal judges from using AI assistance, enabling a productivity response on the adjudication side analogous to the one AI has already produced on the filing side. Another is to more aggressively route simple cases—the cases we show drive the bulk of the post-AI rise—to magistrate judges, lower courts, or specialized triage procedures. These

are not mutually exclusive, and neither is costless. Both are preferable to the alternative, which is a slow accumulation of cases in a system whose throughput is fixed.

Generative AI has lowered the cost of producing passable legal documents. Americans are noticing, and they are walking into federal civil court in much larger numbers to act on it. The courts are not yet visibly buckling under the load—case durations are stable and dispositions look familiar—but the volume of activity inside cases has risen sharply and shows no sign of stopping. Whether the federal civil court system and the broader legal community absorbs this technology as productivity gain or as congestion is being determined in real time.

[jyl: Idiom holds that justice delayed is justice denied. If the rapid diffusion of AI tools has the net effect of overcrowding the federal docket as the cost of filing a case continues to fall, then the increased potential benefits of expanding citizens' access to the courts may be offset by the clamor of those waiting to be heard. It would be a shame if AI tools opened courthouse doors wider only for the line to enter to stretch longer. Then, the scales of justice may tip.]

References

- [1] Administrative Office of the U.S. Courts. *Federal Judicial Caseload Statistics: Judicial Caseload Indicators*. Tech. rep. Administrative Office of the U.S. Courts, 2024. URL: <https://www.uscourts.gov/data-news/reports/statistical-reports/federal-judicial-caseload-statistics/judicial-caseload-indicators-federal-judicial-caseload-statistics-2024>.
- [2] Erik Brynjolfsson, Danielle Li, and Lindsey R. Raymond. “Generative AI at Work”. In: *NBER Working Paper* 31161 (2023).
- [3] Fabrizio Dell’Acqua, Edward McFowland, Ethan R. Mollick, Hila Lifshitz-Assaf, Katherine Kellogg, Saran Rajendran, Lisa Kraye, François Cadelon, and Karim R. Lakhani. “Navigating the Jagged Technological Frontier: Field Experimental Evidence of the Effects of AI on Knowledge Worker Productivity and Quality”. In: *Harvard Business School Working Paper* 24-013 (2023).
- [4] Federal Judicial Center. *Integrated Database: Civil Cases Filed in U.S. District Courts*. Tech. rep. Available at <https://www.fjc.gov/research/idb>. Federal Judicial Center, 2025.
- [5] Mark D. Gough and Emily Taylor Poppe. “(Un)Changing Rates of Pro Se Litigation in Federal Court”. In: *Law & Social Inquiry* 45.3 (2020), pp. 567–596.
- [6] D. James Greiner, Cassandra Wolos Pattanayak, and Jonathan Hennessy. “The Limits of Unbundled Legal Assistance: A Randomized Study in a Massachusetts District Court and Prospects for the Future”. In: *Harvard Law Review* 126.4 (2013), pp. 901–989.
- [7] Neel Guha, Julian Nyarko, Daniel E. Ho, Christopher Ré, Adam Chilton, Aditya Narayana, Alex Chohlas-Wood, Austin Peters, Brandon Waldon, Daniel N. Rockmore, Diego Zambrano, Dmitry Talisman, Enam Hoque, Faiz Surani, Frank Fagan, Galit Sarfaty, Gregory M. Dickinson, Haggai Porat, Jason Hegland, Jessica Wu, Joe Nudell, Joel Niklaus, John Nay, Jonathan H. Choi, Kevin Tobia, Margaret Hagan, Megan Ma, Michael Livermore, Nikon Rasumov-Rahe, Nils Holzenberger, Noam Kolt, Peter Henderson, Sean Rehaag, Sharad Goel, Shang Gao, Spencer Williams, Sunny Gandhi, Tom Zur, Varun Iyer, and Zehua Li. “LegalBench: A Collaboratively Built Benchmark for Measuring Legal Reasoning in Large Language Models”. In: *Advances in Neural*

Information Processing Systems 36 (NeurIPS 2023) Datasets and Benchmarks Track. 2023.

- [8] Gillian K. Hadfield. *Rules for a Flat World: Why Humans Invented Law and How to Reinvent It for a Complex Global Economy*. Oxford University Press, 2016.
- [9] Alex Imas and Brian Jabarian. *Detecting AI-Generated Text*. Working Paper 34223. National Bureau of Economic Research, 2025. URL: <https://www.nber.org/papers/w34223>.
- [10] Daniel Martin Katz, Michael James Bommarito, Shang Gao, and Pablo Arredondo. “GPT-4 Passes the Bar Exam”. In: *Philosophical Transactions of the Royal Society A* 382.2270 (2024). Working paper available at SSRN 4389233 (2023).
- [11] Legal Services Corporation. *The Justice Gap: The Unmet Civil Legal Needs of Low-Income Americans*. Legal Services Corporation, 2022.
- [12] Peter S. Menell and Ryan G. Vacca. “Revisiting and Confronting the Federal Judiciary Capacity “Crisis”: Charting a Path for Federal Judiciary Reform”. In: *California Law Review* 109.4 (2021), pp. 1211–1316.
- [13] Shakked Noy and Whitney Zhang. “Experimental Evidence on the Productivity Effects of Generative Artificial Intelligence”. In: *Science* 381.6654 (2023), pp. 187–192.
- [14] George L. Priest and Benjamin Klein. “The Selection of Disputes for Litigation”. In: *Journal of Legal Studies* 13.1 (1984), pp. 1–55.
- [15] Deborah L. Rhode. “Access to Justice”. In: *Fordham Law Review* 69.5 (2004), pp. 1785–1819.
- [16] Rebecca L. Sandefur. “Access to What?” In: *Daedalus* 148.1 (2019), pp. 49–55.
- [17] Steven Shavell. “Suit, Settlement, and Trial: A Theoretical Analysis Under Alternative Methods for the Allocation of Legal Costs”. In: *Journal of Legal Studies* 11.1 (1982), pp. 55–81.
- [18] Peter Siegelman and Joel Waldfogel. “Toward a Taxonomy of Disputes: New Evidence Through the Prism of the Priest-Klein Model”. In: *Journal of Legal Studies* 28.1 (1999), pp. 101–130.
- [19] Jessica K. Steinberg. “Demand Side Reform in the Poor People’s Court”. In: *Connecticut Law Review* 47.3 (2023), pp. 741–810.
- [20] Joel Waldfogel. “The Selection Hypothesis and the Relationship Between Trial and Plaintiff Victory”. In: *Journal of Political Economy* 103.2 (1995), pp. 229–260.

- [21] Emma Wiles, Zanele Munyikwa, and John J. Horton. “Algorithmic Writing Assistance on Jobseekers’ Resumes Increases Hires”. In: *Management Science* 71.12 (2025). NBER Working Paper 30886.

A. The Vermont outlier: immigration mandamus filings

The District of Vermont is the single state-level outlier in Figure 5. Pro se filings in Vermont rose from roughly 45 cases per year before 2022 to over 1,100 in FY2024. Nearly all of the increase comes from a single NOS category—465, “Other Immigration Actions”—and a single type of claim: writs of mandamus filed against the Director of United States Citizenship and Immigration Services (USCIS). We describe the phenomenon here and show that excluding Vermont from the sample leaves the paper’s results qualitatively unchanged.

Background

The USCIS is the federal agency responsible for processing green-card applications, naturalization petitions, and other immigration benefits. In recent years, processing times for several categories of applications have lengthened substantially, with some applicants waiting two or more years for routine adjudications. A writ of mandamus is a court order compelling a government official to perform a duty owed to the petitioner. When USCIS delays an application beyond what the applicant considers reasonable, the applicant can file a mandamus action in federal court naming the USCIS Director as defendant, asking the court to order USCIS to adjudicate the application.

The District of Vermont became a preferred venue for these filings. It is a small court with a fast docket, and the named defendant—Ur Mendoza Jaddou, Director of USCIS—could be sued there under permissive venue rules. Beginning in 2023, online communities of immigration applicants began sharing detailed instructions for filing mandamus actions *pro se* in the District of Vermont, explicitly leveraging AI tools to draft the filings.

A representative guide

Figure 10 reproduces excerpts from a widely circulated Reddit post offering step-by-step instructions for filing a writ of mandamus without a lawyer. The guide’s first step reads: “I used AI (Microsoft Copilot, to be exact). It’s free and comes with Windows computers. I asked it to write me a writ of mandamus. I knew just AI wasn’t enough, and I know nothing about law, so I did step 2.” Step 2 describes hiring a Fiverr lawyer for \$150 to review and polish the AI-generated draft. Step 3 lists the named defendants: USCIS, the

Attorney General, the USCIS Director, the Secretary of Homeland Security, and the local U.S. Attorney.




Figure 10. Excerpts from a Reddit post (r/USCIS, December 2024) providing step-by-step instructions for filing a *pro se* writ of mandamus using AI and a \$150 Fiverr lawyer review.

The pattern is clear: applicants with no legal training use a large language model to generate an initial draft, pay a small amount for a lawyer to review it, and file the result in a strategically chosen venue. These are legitimate individual cases—each petitioner has a specific pending application and a specific grievance about processing time—but the volume is enabled by the dramatically lower cost of producing the necessary legal documents.

Removing Vermont from the sample

Vermont's contribution to the national *pro se* count is small in absolute terms. Table 1 reports Vermont *pro se* case counts and their effect on the national IDB *pro se* share.

At its peak impact (FY2024), excluding Vermont reduces the national *pro se* share by 0.48 percentage points—from 17.09% to 16.61%. The national rise from the pre-AI baseline of roughly 14% to the FY2025 level of 19.6% is essentially unchanged. Figure 11 plots the national headline shares and counts with and without Vermont; the two series are nearly indistinguishable.

Table 1. Vermont *pro se* cases and impact on national IDB *pro se* share.

FY	Vermont PS		National IDB PS share		
	Cases	× pre-AI	With VT	Without VT	Difference
2019	43	1.0×	14.09%	14.09%	<0.01 pp
2020	47	1.0×	12.60%	12.59%	0.01 pp
2021	45	1.0×	12.48%	12.47%	0.01 pp
2022	51	1.1×	14.09%	14.08%	0.01 pp
2023	307	6.8×	14.94%	14.82%	0.13 pp
2024	1,121	24.9×	17.09%	16.61%	0.48 pp
2025	1,018	22.6×	19.92%	19.55%	0.37 pp




Figure 11. National *pro se* filing counts (left) and share (right) with and without the District of Vermont. The two series are nearly identical. Source: FJC IDB (cv26), non-prisoner non-MDL.

B. Pro se subsampling from CourtListener/RECAP

[avs: Appendix: attempt to restrict the Pangram sample to specifically *pro se* complaints from CourtListener/RECAP. (1) Caveat: RECAP uploads come disproportionately from lawyers and docket-watching services, so the *pro se* subsample is small and selection-biased. (2) Present the results of the subset test.]